

**MINISTRY OF EDUCATION AND TRAINING
CAN THO UNIVERSITY
COLLEGE OF INFORMATION AND COMMUNICATION
TECHNOLOGY
FACULTY OF SOFTWARE ENGINEERING**



**SOFTWARE DEVELOPMENT THESIS PROJECT
COURSE: CT250H**

Thesis title

**DESIGN AND IMPLEMENTATION OF A HUMAN-IN-THE-
LOOP**

**AI-ASSISTED RECRUITMENT AUTOMATION PLATFORM
FOR CV-JD EVALUATION AND RECOMMENDATION IN IT**

Student

Pham Huu Hung - B2203557

Can Tho, August 2026

**MINISTRY OF EDUCATION AND TRAINING
CAN THO UNIVERSITY
COLLEGE OF INFORMATION AND COMMUNICATION
TECHNOLOGY
FACULTY OF SOFTWARE ENGINEERING**



**SOFTWARE DEVELOPMENT THESIS PROJECT
COURSE: CT250H**

Thesis title

**DESIGN AND IMPLEMENTATION OF A HUMAN-IN-THE-
LOOP**

**AI-ASSISTED RECRUITMENT AUTOMATION PLATFORM
FOR CV-JD EVALUATION AND RECOMMENDATION IN IT**

**Supervisor
Ph.D. Nguyen Thanh Khoa**

**Student
Pham Huu Hung - B2203557**

Can Tho, August 2026

ACKNOWLEDGEMENTS

I would like to express my sincere gratitude to Ph.D. Nguyen Thanh Khoa for the guidance, constructive feedback, and continuous support provided throughout this thesis. The supervisor's academic and technical advice played an important role in shaping the research direction and improving the quality of this work.

I am also grateful to the lecturers and staff of the College of Information and Communication Technology, Can Tho University, for providing the academic foundation and learning environment necessary for the completion of this project. My thanks are extended to everyone who contributed feedback during the design, implementation, and evaluation of the CareerFit system.

Finally, I would like to thank my family and friends for their encouragement and support throughout my studies. Although every effort has been made to ensure the accuracy and quality of this thesis, limitations may remain, and constructive comments are sincerely appreciated.

Can Tho, August 2026
Pham Huu Hung

DECLARATION OF ORIGINALITY

I hereby declare that this thesis, titled “Design and Implementation of a Human-in-the-Loop AI-Assisted Recruitment Automation Platform for CV-JD Evaluation and Recommendation in IT,” is my original work completed under the supervision of Ph.D. Nguyen Thanh Khoa. The implementation, experiments, data analysis, and conclusions presented in this report are reported truthfully within the stated scope.

All ideas, publications, standards, software documentation, figures, and other materials obtained from external sources are acknowledged through citations or source notes. I accept full responsibility for the accuracy, academic integrity, and originality of the submitted work.

Can Tho, August 2026
Student

Pham Huu Hung

SUPERVISOR'S COMMENTS

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

Can Tho, August 2026
Supervisor

Ph.D. Nguyen Thanh Khoa

TABLE OF CONTENTS

ACKNOWLEDGEMENTS	i
DECLARATION OF ORIGINALITY	ii
SUPERVISOR'S COMMENTS	iii
TABLE OF CONTENTS.....	iv
LIST OF TABLES.....	vii
LIST OF FIGURES.....	viii
LIST OF ABBREVIATIONS	ix
ABSTRACT	x
INTRODUCTION	1
Problem Statement.....	1
Background and Related Work	2
Research Objectives.....	2
Overall Objective	2
Specific Objectives	2
Research Subjects and Scope.....	3
Research Subject	3
Research Scope	3
Research Methods and Content	4
Research Methods	4
Research Content	4
Key Contributions.....	5
Thesis Structure	5
MAIN CONTENT	6
CHAPTER 1. PROBLEM DESCRIPTION AND REQUIREMENTS	6
1.1 System Analysis Context	6
1.2 Stakeholders and Actors.....	6
1.3 Functional Requirements	8
1.4 Non-Functional Requirements	9
1.5 Use-Case Analysis	10
1.5.1 Candidate Uploads a CV and Receives Matches	10
1.5.2 Candidate Searches, Applies, and Withdraws	11
1.5.3 Recruiter Creates a Job and Reviews Candidates	11
1.5.4 Candidate Feedback Changes Later Ranking.....	12
1.5.5 AutoFit Executes a Policy-Governed Action	12
1.5.6 User Completes an Email Feedback Action.....	13
1.6 Chapter Summary	14
CHAPTER 2. THEORETICAL BACKGROUND AND SOLUTION DESIGN.....	15
2.1 Recruitment Information and CV–Job Description Matching.....	15
2.2 Text Representation and Preprocessing	16
2.2.1 Document Normalization	16
2.2.2 Bag-of-Words Representation.....	17
2.3 TF-IDF and Cosine Similarity	17
2.3.1 Term Frequency and Inverse Document Frequency.....	17
2.3.2 Cosine Similarity.....	17
2.4 Matching, Recommendation, and Ranking.....	18
2.5 Relevance Feedback and the Rocchio Method	19
2.5.1 Relevance Feedback.....	19
2.5.2 Feedback Semantics in CareerFit.....	19
2.6 Evaluation Metrics for Ranked Results	20
2.6.1 Precision@K, Recall@K, and HitRate@K.....	20
2.6.2 Mean Reciprocal Rank.....	21
2.6.3 Discounted Cumulative Gain	21
2.7 Human-in-the-Loop and Explainable Automation.....	21
2.7.1 Levels of Human Involvement.....	21
2.7.2 Explainability and Auditability	22

2.8 Web Architecture, Security, and Audit Foundations	23
2.8.1 Client–Server and REST	23
2.8.2 Authentication, Authorization, and Action Tokens	23
2.8.3 Audit Logging	24
2.9 Related Work and Research Gap	24
2.10 Theoretical Background Summary	25
2.11 System Architecture	25
2.12 Module Design.....	26
2.13 Data Design.....	27
2.13.1 Core Identity and Recruitment Data.....	27
2.13.2 Automation, Communication, and Operational Data	27
2.14 Security, Failure, and Consistency Design	28
2.14.1 Authentication and Authorization	28
2.14.2 Failure Handling	29
2.14.3 Identified Design Risks	29
2.15 Deployment Architecture.....	29
2.16 Chapter Summary	30
CHAPTER 3. SYSTEM IMPLEMENTATION	31
3.1 Implementation Overview	31
3.2 Authentication and Authorization Implementation.....	32
3.2.1 Login and JWT Processing	32
3.2.2 URL Rules and Ownership Checks.....	33
3.3 CV Ingestion and Validation	34
3.3.1 Upload and Manual Creation	34
3.3.2 File Validation and Extraction.....	34
3.4 Text Normalization and TF-IDF Implementation	35
3.4.1 Text Normalization.....	35
3.4.2 Static Corpus and Vector Construction	35
3.5 Matching and Recommendation Implementation.....	36
3.5.1 Scoring Service	36
3.5.2 Matching Orchestration and Persistence	37
3.5.3 Recommendation Service.....	37
3.6 Feedback Learning with Rocchio	38
3.7 Application and Recruiter Workflow	39
3.8 AutoFit and Background Processing.....	40
3.8.1 Async Executor and Scheduler.....	40
3.8.2 Auto-Apply.....	40
3.8.3 Notification Guard and Delivery.....	41
3.9 Email Action and Audit Implementation	41
3.10 Persistence and API Implementation	42
3.11 Frontend Integration.....	43
3.12 Deployment and Observability Implementation.....	47
3.13 Implementation Limitations.....	47
3.14 Chapter Summary	48
CHAPTER 4. TESTING AND EVALUATION	49
4.1 Evaluation Objectives	49
4.2 Evaluation Environment	49
4.3 Dataset and Evaluation Design	50
4.3.1 Controlled Algorithm Dataset	50
4.3.2 Local Runtime Data	51
4.4 Backend Automated Testing	51
4.5 Controlled Algorithm Results	51
4.6 Concurrency Observation in Benchmark Runs.....	52
4.7 Frontend Build Evaluation.....	53
4.8 Browser End-to-End Evaluation.....	53
4.9 Security-Oriented API Checks.....	54
4.10 Runtime Health, Monitoring, and Local Latency	55

4.11 Evaluation Summary by Research Question.....	56
4.12 Threats to Validity.....	56
4.12.1 Construct Validity.....	56
4.12.2 Internal Validity.....	56
4.12.3 External Validity.....	56
4.12.4 Reproducibility.....	57
4.13 Chapter Summary	57
CONCLUSION	58
1. Summary of Results.....	58
2. Achievement of Thesis Objectives.....	59
2.1 Role-Based Recruitment Platform	59
2.2 CV and Job-Description Processing.....	59
2.3 Matching and Recommendation	59
2.4 Feedback Learning.....	59
2.5 Policy-Driven Automation and Email Action	60
2.6 Auditability and Evaluation.....	60
3. Discussion.....	61
3.1 Value of an Interpretable Baseline	61
3.2 Human-in-the-Loop as System Structure.....	61
3.3 Meaning of the Rocchio Improvement	61
3.4 Test Results versus Background Errors.....	62
3.5 Product Positioning	62
4. Limitations.....	62
4.1 Data and Algorithm Limitations.....	62
4.2 Evaluation Limitations	62
4.3 Security and Privacy Limitations	63
4.4 Reliability and Maintainability Limitations	63
5. Future Work	63
6. Closing Remarks.....	64
REFERENCES	65
APPENDICES	67
Appendix A. Requirement-Test Traceability Matrix.....	67
Appendix B. Selected API Contracts	67
Appendix C. Data Model Summary	67
Appendix D. Evaluation Summary	67
Appendix E. UAT and Demonstration Script.....	68
Preparation	68
Execution	68
Acceptance and cleanup.....	68
Appendix F. Local Deployment Instructions	69
Prerequisites	69
Startup procedure	69
Verification.....	69
Shutdown and cleanup	69
Appendix G. Role-Based User Guide.....	70
Guest	70
Candidate	70
Recruiter.....	70
Administrator	70

LIST OF TABLES

Table 1.1. Mapping of objectives, contributions, and evidence	5
Table 1.2. Actors and responsibilities	6
Table 1.3. Consolidated functional requirements	8
Table 1.4. Non-functional requirements and design responses	9
Table 1.5. Use case - Candidate uploads a CV and receives matches	10
Table 1.6. Use case - Candidate searches, applies, and withdraws.....	11
Table 1.7. Use case - Recruiter creates a Job and reviews candidates.....	11
Table 1.8. Use case - Candidate feedback changes later ranking	12
Table 1.9. Use case - AutoFit executes a policy-governed action	12
Table 1.10. Use case - User completes an email feedback action	13
Table 2.1. Positioning of CareerFit against major approach categories	24
Table 2.2. Backend modules and responsibilities	26
Table 2.3. Key integrity constraints	27
Table 2.4. Design risks and required treatment	29
Table 3.1. Main implementation technologies.....	31
Table 3.2. CV processing states	34
Table 3.3. Implemented score interpretation	36
Table 3.4. Feedback handling	39
Table 3.5. Representative API-to-service mapping	43
Table 3.6. Verified implementation limitations	47
Table 4.1. Evaluation environment.....	49
Table 4.2. Backend automated test results.....	51
Table 4.3. Baseline and Rocchio results on the synthetic holdout scenario	51
Table 4.4. Frontend build artifacts	53
Table 4.5. Chromium P0 workflow results.....	53
Table 4.6. API authorization observations.....	54
Table 4.7. Runtime observations	55
Table 4.8. Answers supported by current evidence	56
Table C.1. Consolidated result status.....	58
Table C.2. Objective assessment	60
Table C.3. Principal limitations and impact.....	63

LIST OF FIGURES

Figure 1.1. CareerFit problem context and principal information flows	1
Figure 1.2. Scope boundary of the CareerFit thesis.....	4
Figure 1.3. CareerFit system context.....	6
Figure 1.4. CareerFit use-case overview	8
Figure 1.5. CV ingestion and matching sequence	11
Figure 1.6. Feedback learning and recomputation sequence	12
Figure 1.7. AutoFit decision flow.....	13
Figure 1.8. Secure email-action confirmation and redemption sequence.....	14
Figure 2.1. Distinction between matching, recommendation, and recruitment action	16
Figure 2.2. TF-IDF vectorization and cosine-similarity pipeline	18
Figure 2.3. Rocchio relevance-feedback vector update.....	20
Figure 2.4. Human-in-the-Loop control cycle in CareerFit.....	22
Figure 2.5. CareerFit container and component architecture	26
Figure 2.6. CareerFit logical entity relationships	28
Figure 2.7. Local and containerized deployment topology	30
Figure 3.1. Backend module structure and request flow	32
Figure 3.2. JWT authentication and authorization boundaries	33
Figure 3.3. CV ingestion implementation pipeline.....	35
Figure 3.4. Seed-corpus initialization and TF-IDF construction.....	36
Figure 3.5. Scoring and matching persistence flow.....	38
Figure 3.6. Feedback processing and post-commit learning	39
Figure 3.7. Application and invitation state transitions	40
Figure 3.8. Scheduler and AutoFit execution boundaries.....	41
Figure 3.9. Hashed email-action token and confirm-then-POST flow	42
Figure 3.10. Frontend routes and API data flow.....	44
Screen 3.1. Public Job search	44
Screen 3.2. Candidate matching workspace	45
Screen 3.3. Candidate CV upload interface.....	45
Screen 3.4. Recruiter candidate workspace	46
Screen 3.5. Candidate AutoFit settings.....	46
Screen 3.6. Administrator audit view	47
Figure 4.1. Evaluation environments and evidence sources.....	50
Figure 4.2. Baseline and Rocchio benchmark metrics at $K = 5$	52
Figure 4.3. P0 end-to-end workflow coverage	54
Figure 4.4. Local Job-search latency distribution.....	55

LIST OF ABBREVIATIONS

Abbreviation	Full Term	Meaning / Explanation
AI	Artificial Intelligence	Computer systems designed to perform tasks commonly associated with human intelligence
API	Application Programming Interface	A defined interface through which software components communicate
CV	Curriculum Vitae	A document describing a candidate's qualifications and experience
E2E	End-to-End	Testing of a complete workflow across integrated components
HITL	Human-in-the-Loop	Automation in which humans retain oversight or approval authority
JD	Job Description	A description of a job's responsibilities and requirements
JWT	JSON Web Token	A compact token format used for authentication and authorization
MRR	Mean Reciprocal Rank	A ranking metric based on the position of the first relevant item
nDCG	Normalized Discounted Cumulative Gain	A ranking metric that accounts for relevance and position
NLP	Natural Language Processing	Methods for processing and analyzing human language
OCR	Optical Character Recognition	Extraction of machine-readable text from images or scanned documents
RBAC	Role-Based Access Control	Authorization based on assigned user roles
REST	Representational State Transfer	An architectural style for networked APIs
TF-IDF	Term Frequency–Inverse Document Frequency	A weighting method for representing text terms
UAT	User Acceptance Testing	Validation of whether the system satisfies intended user workflows

ABSTRACT

Background: Recruitment in the information technology sector requires candidates and recruiters to compare many types of information from curricula vitae and job descriptions. Most job portals support posting, searching, and manual applications, while automated screening tools may not clearly explain how a score is calculated or how an automated action is controlled.

Objectives: This thesis designs and implements CareerFit IT AutoPilot, a web-based recruitment platform that combines a job portal, CV-job description matching, personalized job recommendations, policy-based automation, and Human-in-the-Loop interaction in one auditable workflow.

Methods: Recruitment text is normalized and represented with Term Frequency-Inverse Document Frequency vectors. Cosine similarity ranks CV-job pairs and candidate-job recommendations, while Rocchio relevance feedback updates learned representations from explicit feedback. AutoFit checks scores, consent, thresholds, interaction history, quotas, cooldowns, time zones, and quiet hours before choosing an action. CareerFit uses a Spring Boot backend, a React and TypeScript frontend, and PostgreSQL persistence managed through Flyway migrations.

Results: The refreshed backend suite passed 72 of 72 tests, the production frontend build completed successfully, and all 20 Chromium workflow, contract, and resilience tests passed. In the controlled synthetic benchmark, $nDCG@5$ increased from 0.037737 to 0.837737 after Rocchio feedback. These results support a well-tested academic prototype, but they do not establish production readiness or real-world hiring effectiveness.

Keywords: *recruitment automation; CV-job matching; recommendation system; Human-in-the-Loop; Rocchio feedback; explainable automation.*

INTRODUCTION

Problem Statement

IT recruitment requires candidates and recruiters to compare many details, including technical skills, experience, language, location, and working conditions. Candidates must identify suitable vacancies, while recruiters must review CVs and decide which applicants deserve closer attention. As the number of jobs and profiles grows, doing this manually becomes slow and inconsistent [10].

Most job portals support posting, searching, and applying, but they do not fully explain why a job fits a candidate or why one applicant ranks above another. CareerFit therefore treats CV–JD matching and personalized recommendation as related but separate tasks. The first compares a specific CV and job; the second prioritizes jobs from a candidate's broader profile and preferences.

Automation can reduce repetitive work, but a similarity score should not directly trigger an important action. User consent, application state, previous interactions, thresholds, quotas, and special conditions also matter. CareerFit uses a Human-in-the-Loop design so users can review important actions and administrators can trace what happened.

This thesis develops CareerFit IT AutoPilot, an IT-focused web platform that combines job discovery, CV and JD processing, matching, recommendation, Rocchio feedback, AutoFit policies, email actions, and audit records. The aim is not to replace recruitment judgment, but to provide understandable evidence and controlled assistance throughout the workflow.

CareerFit problem context



CareerFit IT AutoPilot — thesis diagram

Figure 1.1. CareerFit problem context and principal information flows

Background and Related Work

The first motivation is to reduce fragmentation. Candidates often move between job search, CV management, application history, and email, while recruiters switch between vacancy management, applicant review, candidate discovery, and reporting. CareerFit brings these activities into one role-based system and uses email as an additional action channel.

The second motivation is explainability. TF-IDF and cosine similarity do not capture meaning as deeply as modern embedding models, but their terms and weights can be inspected. This makes them suitable for an academic baseline in which the scoring process, limitations, and effect of feedback must be demonstrated clearly.

The third motivation is controlled adaptation and automation. AutoFit places policy checks between a score and an action, considering consent, thresholds, previous interactions, cooldown, quota, time zone, and quiet hours. Rocchio feedback then provides a transparent way to adjust later rankings from explicit positive and negative judgments without overwriting the original representation.

Related work includes classical lexical retrieval, learned resume-job representations, explainable recommendation, and human-centered governance. Lexical methods provide inspectable weights and matching reasons, while recent dense approaches improve semantic representation but require stronger data and evaluation assumptions [6]-[8]. Research on algorithmic hiring and explainable recommendation also shows that model performance, fairness, user control, and explanation quality must be evaluated separately [9], [10], [14]. CareerFit addresses a narrower integration gap by combining portal workflows, explainable scoring, feedback learning, policy-controlled actions, and audit records in one academic prototype.

Research Objectives

Overall Objective

The overall objective of this thesis is to design, implement, and evaluate a web-based Human-in-the-Loop recruitment automation platform for the IT domain. The platform is intended to support CV–job description evaluation, personalized job recommendation, feedback-based adaptation, and policy-driven actions while preserving user control, explainability, security, and auditability.

Specific Objectives

To achieve the overall objective, the thesis defines the following specific objectives:

- Design a role-based web platform for guests, candidates, recruiters, and administrators, with appropriate public and authenticated workflows.
- Provide candidate profile and CV management, including document upload, form-based entry, text extraction, OCR fallback, validation, and quality feedback.

- Implement separate CV–JD matching and candidate-profile recommendation workflows using a shared text-normalization, TF-IDF, and cosine-similarity pipeline.
- Present normalized scores, relevance labels, matching reasons, validation signals, and stable ranking behavior to support interpretation by candidates and recruiters.
- Integrate explicit user feedback and apply the Rocchio relevance-feedback method to update learned representations and recompute rankings reproducibly.
- Implement AutoFit policies that convert matching results into controlled outcomes according to consent, thresholds, interaction state, quota, cooldown, time-zone, and quiet-hour rules.
- Support expiring email-action links with hashed token storage, a non-mutating confirmation page, POST execution, action-state tracking, and audit logging.
- Evaluate the algorithmic behavior and system workflows using reproducible experiments, automated tests, scripted acceptance scenarios, and security and reliability checks appropriate to the project scope.

Research Subjects and Scope

Research Subject

The research subject is the IT recruitment workflow supported by a web platform, including CV and Job-description text, lexical ranking and recommendation, explicit feedback signals, policy-controlled automation, and auditable user and system actions.

Research Scope

This thesis focuses on an academic prototype for recruitment in the information technology domain. Its functional scope includes a public job portal; authenticated candidate, recruiter, and administrator workspaces; CV and profile management; job and employer management; CV–JD matching; personalized job recommendation; application workflows; explicit relevance feedback; AutoFit policy configuration; notifications and actionable email; analytics; and audit records.


The data-processing scope includes extraction of text from supported CV documents, OCR fallback for image-based content where the configured runtime permits it, hard and soft validation, bilingual-oriented text normalization, TF-IDF vectorization, cosine-similarity scoring, score normalization, relevance labeling, and Rocchio-based feedback updates. PostgreSQL is used as the primary data store, and database evolution is managed through Flyway migrations. The implemented system follows a client–server architecture with a React and TypeScript frontend and a Spring Boot backend exposing REST APIs.

The thesis does not attempt to implement a complete enterprise applicant tracking system. Interview scheduling, offer management, payroll, and full employee lifecycle management are outside the core scope. The system does not automatically submit applications to external third-party job boards, and it does not use a large language

model to make unrestricted recruitment decisions. Distributed microservices, large-scale message brokers, and semantic embedding models may be discussed as future extensions but must not be presented as implemented core components unless they are added and verified before the final report is submitted.

The evaluation scope is limited. Controlled or synthetic datasets can show whether Rocchio changes rankings in the expected direction, but they cannot prove recruitment quality in a real production setting. The source and purpose of scraped, seeded, or synthetic data must be stated. Claims about usability, performance, security, or production readiness are limited to the test environment, participants, procedures, and evidence available during the final evaluation.

Thesis scope boundary



CareerFit IT AutoPilot — thesis diagram

Figure 1.2. Scope boundary of the CareerFit thesis

Research Methods and Content

Research Methods

This thesis follows an applied software-engineering research approach. The work combines requirement analysis, architecture and data design, incremental implementation, a controlled algorithm experiment, automated backend and browser testing, runtime observation, and threats-to-validity analysis. Implementation claims are checked against source code, migrations, API contracts, logs, test reports, and generated evaluation artifacts.

Research Content

The research models CareerFit as a Perception-Decision-Action-Learning-Audit workflow. It first identifies recruitment roles, requirements, and control boundaries; then designs and implements the data, matching, recommendation, feedback, automation, email action, and audit modules. Finally, the study evaluates algorithm

behavior, software correctness, browser workflows, authorization, runtime health, and local latency within the stated experimental scope.

Key Contributions

The principal contribution of this thesis is the design and implementation of an integrated, controlled recruitment workflow rather than the invention of a new machine-learning model. The contributions are summarized as follows:

- An integrated recruitment platform that combines a job portal, CV–JD matching, personalized recommendation, application management, and role-specific operational workspaces.
- A transparent text-scoring pipeline using TF-IDF and cosine similarity, with normalized results, labels, reasons, validation signals, and explicit separation between matching and recommendation objectives.
- A feedback-learning workflow based on the Rocchio method, designed to update learned representations from explicit feedback and to support reproducible recomputation rather than uncontrolled cumulative modification.
- A policy-driven AutoFit layer that separates relevance scoring from business actions and applies Human-in-the-Loop controls, user consent, operational constraints, and previous interaction states.
- An actionable email mechanism with expiring, hashed tokens, a non-mutating GET confirmation page, POST redemption, and action-state tracking.
- A traceable evaluation structure that distinguishes algorithm experiments from functional, integration, acceptance, performance, reliability, and security evidence, while explicitly documenting threats to validity.

Thesis objective	Implemented contribution	Evaluation evidence
Analyze IT recruitment workflows	Role-specific requirements and Human-in-the-Loop boundaries	Use cases and traceability matrix
Implement CV–JD matching	TF-IDF, cosine scoring, explanations, and persisted Matchings	Algorithm benchmark and backend tests
Learn from feedback	Rocchio updates followed by scheduled recomputation	Baseline-versus-learned ranking metrics
Support controlled automation	AutoFit policy, quota, cooldown, notification, and audit controls	Security tests and P0 E2E flows
Deliver an integrated platform	Spring Boot, React/TypeScript, PostgreSQL, and Flyway	Build, runtime health, and browser evidence

Table 1.1. Mapping of objectives, contributions, and evidence

Thesis Structure

The report contains an Introduction, four chapters, and a Conclusion. Chapter 1 defines the problem, stakeholders, requirements, and use cases. Chapter 2 presents the theoretical background and solution design. Chapter 3 explains the implementation. Chapter 4 presents the evaluation method, evidence, and threats to validity. The Conclusion discusses achievements, limitations, and future work.

MAIN CONTENT


CHAPTER 1. PROBLEM DESCRIPTION AND REQUIREMENTS

1.1 System Analysis Context

CareerFit is designed as a role-based system for IT recruitment. Its boundary includes the React frontend, Spring Boot backend, PostgreSQL database, CV storage, background scheduler, and optional email provider. External job boards, interview scheduling, offers, payroll, and unrestricted automated hiring remain outside the implemented scope.

The central design distinction is between evidence, business state, and action. A matching record contains evidence about the similarity between one CV and one job. An application record represents a candidate's participation in a recruitment process. An automation policy represents consent and operational preferences. These concepts must remain separate: a high similarity score must not silently become an application, and an application status must not be interpreted as a relevance label.

System context



CareerFit IT AutoPilot — thesis diagram

Figure 1.3. CareerFit system context

1.2 Stakeholders and Actors

CareerFit serves four interactive roles and two supporting actors. Guest is a frontend state rather than a persisted account role. Candidate, Recruiter, and Administrator are persisted roles in user_account. Background processing and the mail provider are non-human actors that execute scheduled work and deliver messages.


Table 1.2. Actors and responsibilities

Actor	Primary responsibilities	Access boundary
-------	--------------------------	-----------------

Actor	Primary responsibilities	Access boundary
Guest	Browse public jobs, search suggestions, job details, featured employers, and public market analytics	Public GET operations only; authentication is required for personalized data or application actions
Candidate	Maintain profile, CVs, and portfolio; view matching and recommendation results; apply or withdraw; provide feedback; configure personal automation	Candidate-owned resources and candidate APIs
Recruiter	Maintain employer profile and jobs; inspect applicants and discovered candidates; invite candidates; update application status; view recruiter analytics	Recruiter APIs and jobs owned by the authenticated recruiter
Administrator	Monitor system summaries, users, jobs, email actions/tokens, and audit records; moderate jobs and trigger controlled maintenance operations	Administrative APIs only
Background scheduler	Recompute stale matches, send digests and notifications, clean expired email actions, and execute auto-apply scans	Internal service calls under system control
Mail provider	Deliver passwordless or notification email generated by the backend	SMTP boundary; it does not own CareerFit business state

Stakeholder needs sometimes conflict. Candidates want relevant opportunities, privacy, and control over automatic applications. Recruiters want efficient prioritization without losing access to applicants who do not rank highly. Administrators need visibility and recoverability without unrestricted exposure of CV content. The design therefore combines role authorization, ownership checks, policy settings, validation, explicit statuses, and audit records rather than relying on a single ranking score.

Role-oriented use cases



CareerFit IT AutoPilot — thesis diagram

Figure 1.4. CareerFit use-case overview

1.3 Functional Requirements

The detailed Software Requirements Specification contains individual requirement identifiers. For the thesis, they are consolidated into functional groups that map to implemented modules and observable workflows.

Table 1.3. Consolidated functional requirements

Group	Required capabilities	Main actor
Authentication and account	Registration, password login, passwordless verification, current-account lookup, role enforcement, account settings	Candidate, Recruiter, Administrator
Public job portal	Search, suggestions, filters, job details, featured employers, employer details, and public analytics	Guest and authenticated users
Candidate profile and CV	Profile and portfolio maintenance, multiple CVs, default CV selection, document upload, manual creation, status tracking, validation, extraction, and OCR fallback	Candidate
Job and employer management	Employer profile management; job creation, update, status change, deletion, search, and CSV export	Recruiter
Matching and recommendation	CV–JD scoring, labels, reasons, candidate match cards, job ranking, profile-based recommendations, and similar jobs	Candidate and Recruiter
Application workflow	Manual apply, candidate application history, withdrawal, recruiter applicant view, invitation, and status update	Candidate and Recruiter

Group	Required capabilities	Main actor
Feedback learning	Record explicit feedback from web or email, update learned representations with Rocchio, and mark affected matches for recomputation	Candidate
Automation and notification	Policy retrieval/update, pause/resume, run-now, scheduled auto-apply, digest, high-match email, quota, cooldown, and delivery logging	Authenticated user and Scheduler
Analytics and administration	Market, candidate, recruiter, and system summaries; event tracking; user/job moderation; audit and email-token monitoring	Candidate, Recruiter, Administrator

The implementation exposes these groups through REST-oriented endpoints. Examples include `/api/jobs/search`, `/api/cv/upload`, `/api/matches/me/cards`, `/api/recommendations/jobs`, `/api/applications`, `/api/recruiter/jobs/{jobId}/candidates`, `/api/automation/policy`, and `/api/admin/audit-logs`. Chapter 3 describes endpoint behavior in detail. This chapter uses endpoint groups only to establish functional boundaries and traceability.

1.4 Non-Functional Requirements

Non-functional requirements constrain how the workflows must operate. They are expressed as design targets; Chapter 4 must distinguish targets from properties verified by measurement.

Table 1.4. Non-functional requirements and design responses

Quality attribute	Requirement	Design response
Security	Protected resources must require authenticated and authorized access	Stateless Spring Security filter chain, JWT processing, role rules, ownership checks, BCrypt password hashing, CORS allow-list, CSP, and HSTS configuration
Privacy	CV and account data must not be public or unnecessarily logged	Candidate-scoped APIs, local protected storage boundary, selective DTOs, and audit metadata rather than full document content
Reliability	Repeated or failed operations should not create inconsistent domain state	Database constraints, transactions, unique keys, optimistic version columns, idempotency checks, explicit statuses, and retry-aware background work
Performance	Public queries and ranking views should remain responsive for the academic dataset	Database indexes, pagination, stored vectors, background scoring, and bounded result sets
Maintainability	Domain changes should remain localized and database evolution reproducible	Modular package structure, controller–service–repository separation, DTOs, configuration properties, Flyway migrations, and automated tests